

Until Wednesday, Jan. 7th, 2026, 11:59 pm CET, solutions to the following exercises must be submitted as one zip-file named `ML25-ex4-group<your-group-number>.zip` via Moodle:
[1](#), [2](#), [3](#), [4](#), [5](#), [7a,b](#), and [8](#).

Exercise 1 : Gradient Descent (1+1+1=3 Points)

- (a) Name one difference between the perceptron training rule and the gradient descent method.
- (b) What are the main requirements for the hypothesis space and error function for gradient descent to be successfully applied?
- (c) Name the key difference between the algorithm of batch gradient descent (BGD) and the algorithm of incremental gradient descent (IGD).

Exercise 2 : Perceptron Learning (2 Points)

Consider two perceptrons, $y_0()$ and $y_1()$, both defined by the function $\text{heaviside}(\sum_{j=0}^p w_j x_j)$. Both perceptrons have identical weights except for the bias term: $w_0 = 0$ for $y_0()$ and $w_0 = 1$ for $y_1()$. Determine if one of the perceptrons, $y_0()$ or $y_1()$, is *more general* than the other (as defined in the lecture units on concept learning). If one is more general, specify which one and explain your answer.

Exercise 3 : Backpropagation (2 Points)

In the early slides [ML:IV-9](#), the Perceptron uses the Heaviside step function. In the MLP section [ML:IV-66](#), this is replaced by the Sigmoid function $\sigma(z)$. Explain the primary mathematical reason for this switch in the context of the Backpropagation algorithm.

Exercise 4 : Perceptron Learning (1+1+2+2+1=6 Points)

In this exercise, you design a single perceptron with two inputs x_1 and x_2 . This perceptron shall implement the boolean formula $A \wedge \neg B$ with a suitable function $y(x_1, x_2)$. Use the values 0 for *false* and 1 for *true*.

- (a) Draw all possible examples and a suitable decision boundary in a coordinate system.
- (b) Draw the graph of the perceptron. The schematic must include x_1 , x_2 , and all model weights.
- (c) Manually determine the weights $\mathbf{w} = (w_0, w_1, w_2)$ for the decision boundary you drew in (a).
- (d) Now determine \mathbf{w} using the perceptron training algorithm (PT). Use a learning rate η of 0.3 and initialize the weights with $w_0 = -0.5$ and $w_1 = w_2 = 0.5$. Instead of selecting examples randomly, use the following examples in the given order (stop after those four examples):

x_1	x_2	c
0	0	0
0	1	0
1	0	1
1	1	0

Draw the decision boundary after every weight update into the coordinate system of (a).

- (e) Briefly describe one effect of changing the learning rate η on the learning progress.

Exercise 5 : Perceptron Learning (1 Points)

Explain why a single perceptron cannot learn the Boolean XOR function $A \text{ } XOR \text{ } B$. Support your answer with an illustration.

Exercise 6 : Multilayer Perceptron (0 Points)

Consider the minimum multilayer perceptron that can handle the XOR problem (slide ML:IV-59), but with an activation function $\mathbf{f}(\mathbf{z}) = (f(z_1), \dots, f(z_d))^T$ instead of *Heaviside()*. This is a two-layer perceptron with $p = 2$ attributes, a hidden layer with $l = 2$ units, and an output layer with $k = 1$ units.

Therefore, as per slide ML:IV-82:

$$y(\mathbf{x}) = \mathbf{f} \left(W^o \begin{pmatrix} 1 \\ \mathbf{f}(W^h \mathbf{x}) \end{pmatrix} \right) = \mathbf{f} \left((w_{10}^o, w_{11}^o, w_{12}^o) \begin{pmatrix} 1 \\ \mathbf{f} \left(\begin{pmatrix} w_{10}^h & w_{11}^h & w_{12}^h \\ w_{20}^h & w_{21}^h & w_{22}^h \end{pmatrix} \begin{pmatrix} 1 \\ x_1 \\ x_2 \end{pmatrix} \right) \end{pmatrix} \right)$$

Confirm for yourself that thus

$$y(\mathbf{x}) = f(w_{10}^o + w_{11}^o \cdot f(w_{10}^h + w_{11}^h \cdot x_1 + w_{12}^h \cdot x_2) + w_{12}^o \cdot f(w_{20}^h + w_{21}^h \cdot x_1 + w_{22}^h \cdot x_2))$$

Show that this network is not able to learn a solution to the XOR problem if the scalar activation function $f()$ is the identity function: $f(z) = z$.

Exercise 7 : Parameters of the Multilayer Perceptrons (2+1+0+0=3 Points)

In this exercise, you analyze the number of weights (parameters) of multilayer perceptrons. We use the notation from the lecture (e.g., slide [ML:IV-104](#)), where multilayer perceptrons have d layers, p attributes, hidden layer i with l_i units, and an output layer with k units.

- (a) Let $d = 4$, $p = 6$, $l_1 = 4$, $l_2 = 4$, $l_3 = 2$, and $k = 3$. Draw the graph of the multilayer perceptron.
- (b) Calculate the number of weights in the multilayer perceptron of (a).
- (c) Calculate the number of weights in the multilayer perceptron of (a) but with each l_i doubled, i.e., $l_1 = 8$, $l_2 = 8$, $l_3 = 4$. Has the number of weights doubled as well?
- (d) Let $f(p, l_1, \dots, l_{d-1}, k)$ be a function that computes the number of weights in the general case. Write down an expression for f .

Exercise 8 : **P** Classification with Neural Networks (1+1+1+1+1=6 Points)

In this exercise, you will implement a two-layer perceptron for predicting whether a given text was written by a human or generated by a language model. Download and use these files from Moodle (the `tsv` files are the same as in the last sheet):

- `features-train.tsv`: Feature vectors for each example in the training set.
- `features-test.tsv`: Feature vectors for each example in the test set.
- `labels-train.tsv`: Labels for each example in the training set indicating the class `is_human` ($C = \{\text{True}, \text{False}\}$)
- `programming_exercise_neural_networks.py`: Template containing function stubs for the exercise. Use it with the files above as follows:

```
python3 programming_exercise_neural_networks.py
    features-train.tsv labels-train.tsv  features-test.tsv
    predictions-test.tsv
```
- `requirements.txt`: Requirements file for the template; can be used to install dependencies.

- Implement a function `encode_class_values` to encode class values as vectors $\mathbf{c} \in \{0, 1\}^k$ (see slide [ML:IV-100](#)).
- Implement a function `predict_probabilities` that, given W_h and W_o , predicts the class probabilities for each example of a dataset.¹
- Implement a function `predict` that, given W_h and W_o , predicts the class for each example of a dataset (as the one with the highest probability).
- Implement a function `train_multilayer_perceptron` that fits a multilayer perceptron using the IGD Algorithm (slide [ML:IV-100](#)). Like in the last exercise sheet, make sure to return the misclassification rate on both training and validation set for plotting, but this time also return the weights after each iteration (for task (e)).¹
- Select the model (iteration) that achieved the best misclassification rate on the validation set and use it to predict the label for each example in the test set (`features-test.tsv`). Write the predictions as one column to a file `predictions-test.tsv` and submit that along with your other solutions.
- Which lines of your code would you need to change from two classes to three? Mark each of them with the comment “# change for 3 classes”.

If you would like to improve your model, here are some suggestions:

- Try experimenting with adjusting various hyperparameters of the model (e.g. learning rate, number of hidden units, etc).
- Try to decrease the learning rate with each iteration.
- Try to implement one of the feature scaling methods explained in [this video by Andrew Ng](#).
- Implement k -fold cross-validation instead of using a single hold-out in order to get a more robust estimate of your model’s performance on the unseen data.
- Expand your implementation to support perceptrons with more than two layers. Hint: the parameter `l` can be replaced by a list of layer sizes; the weight matrices for the hidden layers can also be stored in a list.

¹ Hint: You can find it in the [code linked on the lecture slides](#).